




# A Survey of Monte Carlo Tree Search Methods




# Introduction

- Consider a game-playing agent
  - How to find good/best next move?
- Game represented as a game tree
  - Nodes represent (subset of) states
  - Still too large to check all actions
- Idea: Combine **sampling** and **tentative rewards**
  - refine “view” of the current game tree iteratively
- Don't overthink the term “game”
  - Generally applicable search algorithm




# Basics: Monte Carlo Tree Search


(1) Selection




(2) Expansion



(3) Simulation




(4) Backpropagation



# Basics: Setting


- $s(v)$  – state  $s$  represented by node  $v$
- $a(v)$  – incoming action  $a$  of node  $v$
- $Q(v)$  – total simulation reward of  $v$
- $N(v)$  – number of times  $v$  has been visited
- Add. data structures possible
  - e.g.  $N(a)$



# Basics: Monte Carlo Tree Search

---


- First Step: Selection
- Controlled by the **Tree Policy**
- Transfer: Multi-Armed Bandit Problem
  - T-times: pick some action
  - collect associated reward
  - maximize total reward
- Later: UCT



# Basics: Monte Carlo Tree Search

---


- Second Step: Expansion
- Controlled by **Tree Policy**
- Transfer: Multi-Armed Bandit Problem
  - T-times: pick some action
  - collect associated reward
  - maximize total reward
- Later: UCT



# Basics: Monte Carlo Tree Search

---


- Third Step: Simulation
- Play out a game
  - i.e. execute actions until a terminal state is reached
  - evaluate result
- Controlled by **Default Policy**
- Details to consider:
  - speed vs. quality
  - complexity of evaluation function



# Basics: Monte Carlo Tree Search

---

- Fourth Step: Backpropagation
- Update values of
  - $Q(v)$
  - $N(v)$
  - ...
- Can be modified by heuristics
  - later - AMAF





# Basics: Monte Carlo Tree Search

---

- About **iterations**:
- While the computation budget is not exhausted
  - 1) execute the four steps as many times as possible
  - 2) Increase tree size and number of evaluated states
  - 3) Finally, pick the best one
- Time-constraints are crucial
  - real-time, turn-based games

# Basics: UCT

---

- “Upper Confidence Bound for Trees”
- Exploitation: Follow discovered path deemed valuable
- Exploration: Discover new paths

“Exploitation”



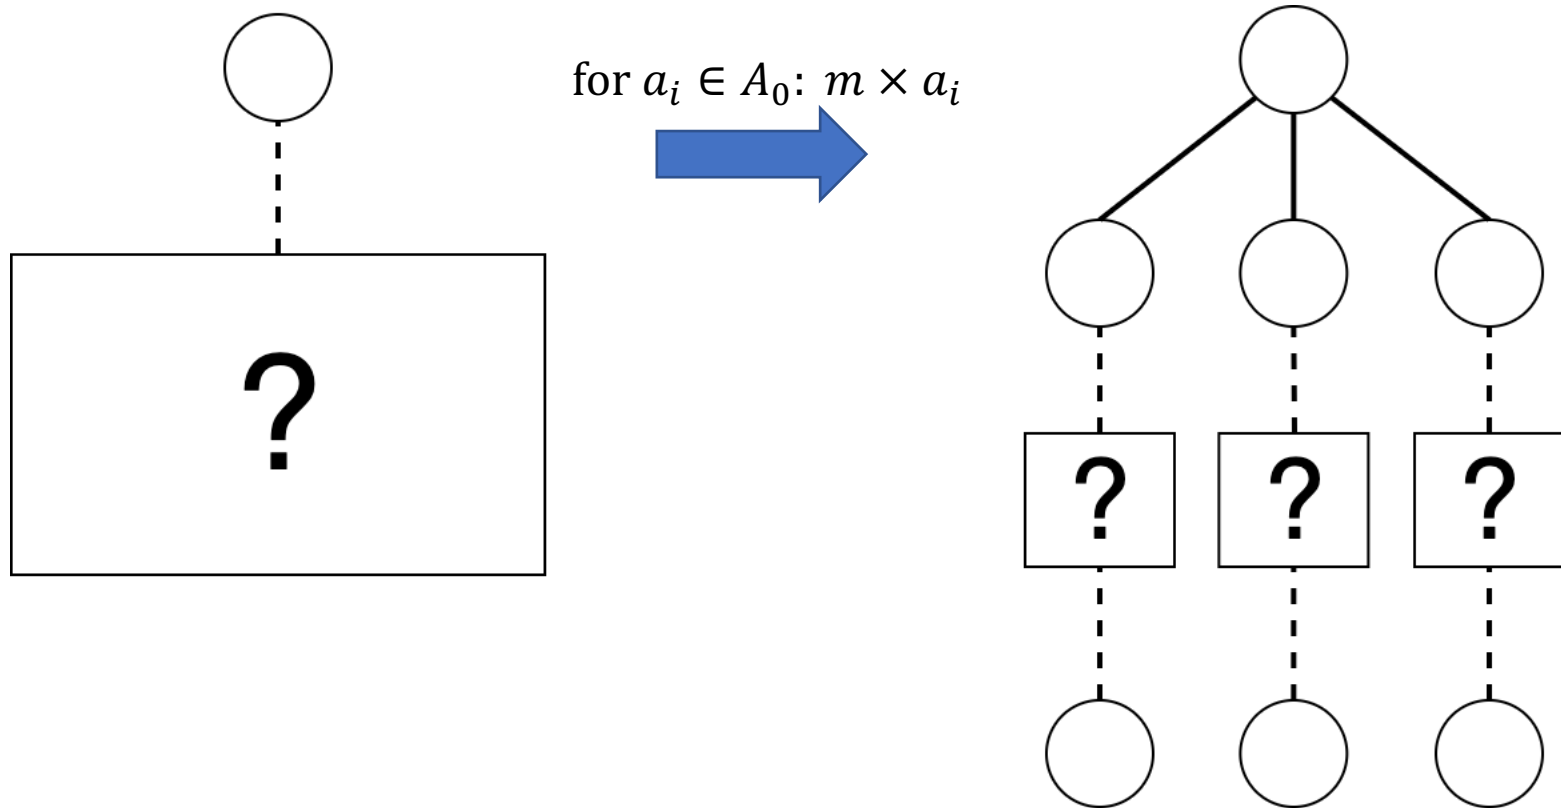
$$UCT(v) = \frac{Q(v')}{N(v')} + c_p \cdot \sqrt{\frac{2 \log N(v)}{N(v')}}}$$



“Exploration”


# Heuristics: BFS-Tree Initialization

- Before MCTS: execute each available root-action
  - Information gained can guide algorithm



# Heuristics: Loss Avoidance

- Ignore negative results when encountering new nodes
- If unvisited node represents a loss, generate neighbors
- Backpropagate only best-case scenario
- Similar (inverse) approach for overly optimistic evaluations



# Heuristics: Novelty-based Pruning


---

- Problem: Many similar states and redundant paths
- Solution: Use a novelty measure  $\text{nov}(v)$ , all nodes above a threshold are pruned
- Step 1: For node  $v$ , define a neighborhood of other nodes
- Step 2: Compare nodes in the neighborhood to  $v$ , giving us their **similarity** or **novelty**
- Step 3: Prune  $v$  if it is too similar to any other nodes



# Heuristics: Novelty-based Pruning

- Step 1: the neighborhood  $N(v)$  is the union of
  - the *left* siblings of  $v$
  - the parent of  $v$ :  $p(v)$
  - *all* siblings of  $p(v)$
  - the neighborhood of  $p(v)$




# Heuristics: Novelty-based Pruning

---

- Step 2: Compare nodes in the neighborhood to  $v$ , giving us their similarity or novelty
- (Somewhat) more precisely:  $nov(v, N(v))$  is
  - the size of the smallest set of predicates that are true for  $v$  and false for all other nodes in  $N(v)$
- For example: if  $save(v)$  and  $\neg save(v')$  where  $v' \in N(v)$ 
  - we have  $nov(v, N(v)) = 1$
- Step 3: Prune all nodes with novelty larger than  $T > 1$

# Heuristics: Knowledge-based Evaluations

- Terminal state  $s_T$  has same value as starting point  $s_0$ 
  - Is it an improvement?
  - classify nearby objects: (state  $i$ , weight  $w_i$ )
  - A\*-Algorithm: distance to objects  $d(i)$
  - Add. value of  $\sum w_i \cdot (d_0(i) - d_{T(i)})$





# Heuristics: All Moves As First

---

- RAVE – Combine UCT and AMAF

- $AMAF(v) = \beta(v) \cdot Q(v) + (1 - \beta(v)) \cdot AMAF(v')$

- where  $\beta(v) = \sqrt{\frac{K}{3 \cdot N(v) + K}}$

ancestor of v: which one?





- Equivalence Parameter K: number of simulations where UCT and AMAF have equal weight
- Generalization: GRAVE – modify ancestor selection procedure



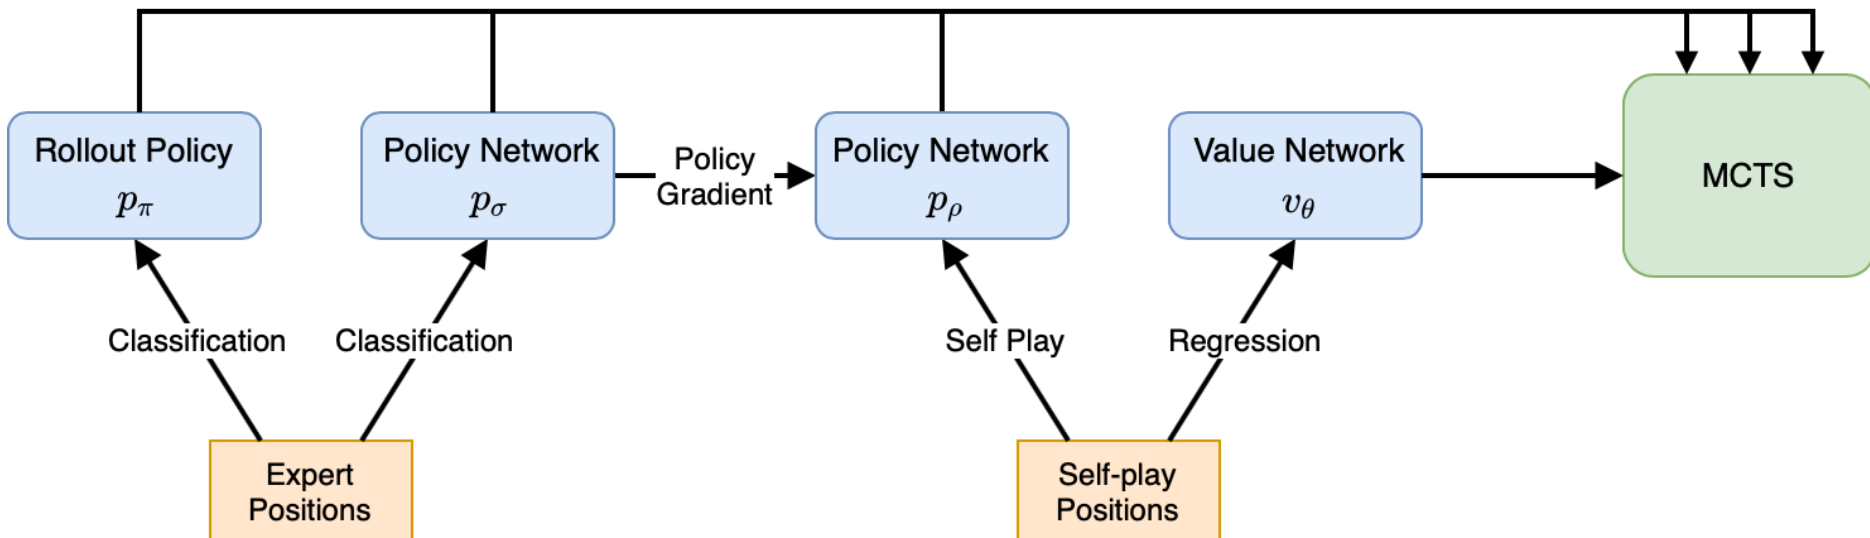
# Extension: Tree Compression

- Control size and width of tree when there are many paths
  - Binary-search
- Inspired by: Deep Architect
  - (Hyper-)parameters for neural network architectures



# Application: AlphaGo

- ML-System that beat some of the world's greatest players in the game of Go
- MCTS used to train a neural network
  - **Lookahead search** combines policy networks and value networks



# Conclusion

---

- MCTS – Search algorithm that combines
  - sampling
  - and simulations
- Addresses Exploration-Exploitation tradeoff
- Works well out of the box but can be extended and tailored to specifics of current domain
  - Heuristics
  - Transfer techniques from statistics and probability theory
- Most famous application: games
  - AlphaGo